

Movie Revenue Prediction

Project overview

The client is a movie studio and they need to be able to predict movie revenue in order to greenlight the project and assign a budget to it. Most of the data is comprised of categorical variables. While the budget for the movie is known in the dataset it is often an unknown variable during the greenlighting process.

Prediction Basis

Since predicting a movie's revenue in itself is a challenge, additionally the budget amount remains elusive makes it a difficult business problem as there is no base-line to establish the prediction.

I- Executive Summary

Analyze available data to predict a movie revenue and finally assign a budget during green-lighting process.

II- Business Understanding

Define Organization

Client is a movie studio in the business of producing, sponsoring and financing movies.

Intended Stakeholders of Data

Movie producers, finance & budget managers.

Define Business Objectives

based on the input parameters, a movie's revenue is to be predicted.

Background

Business Objectives

- Which variables can help in predicting revenue figures?
- determine which parameters had the most effect on a movie revenue?
- Limiting the problem to predicting just the revenue amounts.
- Evaluate and cross-validate the revenue figures.

Business Success Criteria

Predicting the movie revenue as accurately as possible, and making the **REVENUE** as the **TARGET VARIABLE** or **OUTCOME**. Since the revenue is a whole-number, a Regression will be developed.

Assumptions, and Constraints

Assumptions: data is accurate and reliable.

Terminology - Code book - Data Dictionary

A lead is a person who has indicated interest in your company's product or service in some way, shape, or form.

- title - title of the movie
- tagline - few words for movie presentation
- revenue - revenue generated by the movie
- budget - planned expenditure
- genres - categorical group of the movie
- homepage - movie promotional website
- id - movie id
- keywords - tags associated with the movie
- original_language - original movie language
- overview - movie synopsis
- production_companies - sponsoring and producing companies
- production_countries - locations of the movie made
- release_date - movie available for viewing date.
- runtime - movie duration
- spoken_languages - spoken languages in the movie
- status - movie status for viewing

Project Plan

gets updated as per the succeeding stages.

modelling planned to use: Multiple Linear Regression, SVM,

Understanding the Data: Data quality report

```
In [4]: mvrevenue.iloc[0]
```

```
Out[4]: title                                Avatar
tagline                                Enter the World of Pandora.
genres                                [{"id": 28, "name": "Action"}, {"id": 12, "nam...
homepage                                http://www.avatarmovie.com/
id                                19995
keywords                                [{"id": 1463, "name": "culture clash"}, {"id":...
original_language                                en
overview                                In the 22nd century, a paraplegic Marine is di...
production_companies                                [{"name": "Ingenious Film Partners", "id": 289...
production_countries                                [{"iso_3166_1": "US", "name": "United States o...
release_date                                12/10/2009
runtime                                162
spoken_languages                                [{"iso_639_1": "en", "name": "English"}, {"iso...
status                                Released
budget                                237000000
revenue                                2787965087
Name: 0, dtype: object
```

```
In [5]: # check the column type
element = mvrevenue.iloc[0]['genres']
print(type(element))
print(element)
```

```
<class 'str'>
[{"id": 28, "name": "Action"}, {"id": 12, "name": "Adventure"}, {"id": 14, "nam
e": "Fantasy"}, {"id": 878, "name": "Science Fiction"}]
```

Since some cols looks like in dictionary format but in actual its string format

This can be handled as first convert the type string into dictionary format and then extract the values.

```
In [4]: # Create a function to convert string into dict

dict_columns = [ 'genres', 'production_companies', 'production_countries', 'spoken_l
anguages', 'keywords']

def text_to_dict(df):
    for column in dict_columns:
        df[column] = df[column].apply(lambda x: {} if pd.isna(x) else ast.literal_e
val(x) )
    return df
```

```
In [5]: # Convert string into dict
con_data = text_to_dict(mvrevenue)
con_data.head(1)
```

Out [5]:

	title	tagline	genres	homepage	id	keywords	original_language	overview	produc
0	Avatar	Enter the World of Pandora.	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 878, 'name': 'Science Fiction'}]]	http://www.avatarmovie.com/	19995	[[{'id': 1463, 'name': 'culture clash'}, {'id': 1464, 'name': 'marine'}]]	en	In the 22nd century, a paraplegic Marine is di...	[[{'name': 'P...

```
In [8]: # checking for dict format
element = con_data.iloc[0]['genres']
print(type(element))
print(element)

<class 'list'>
[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 878, 'name': 'Science Fiction'}]
```

Initial Data Collection Report

Describing Data at High Level

```
In [9]: # rowsXcolumns format
con_data.shape

# missing values
con_data.info()
```

Out [9]: (4803, 16)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 16 columns):
title                4803 non-null object
tagline              3959 non-null object
genres               4803 non-null object
homepage             1712 non-null object
id                   4803 non-null int64
keywords             4803 non-null object
original_language    4803 non-null object
overview             4800 non-null object
production_companies 4803 non-null object
production_countries 4803 non-null object
release_date         4802 non-null object
runtime              4801 non-null float64
spoken_languages     4803 non-null object
status               4803 non-null object
budget               4803 non-null int64
revenue              4803 non-null int64
dtypes: float64(1), int64(3), object(12)
memory usage: 600.5+ KB
```

```
In [6]: # dataframe bckup copy
#conv_data
conv_data = con_data.copy()
```

```
In [7]: # Checking the null values
con_data.isnull().sum()
```

```
Out[7]: title                0
tagline                    844
genres                     0
homepage                  3091
id                         0
keywords                  0
original_language         0
overview                   3
production_companies      0
production_countries      0
release_date              1
runtime                   2
spoken_languages          0
status                    0
budget                    0
revenue                   0
dtype: int64
```

Since homepage, tagline has the most missing values and don't add any weight, since these are text columns, in building prediction model, these can be dropped.

```
In [8]: con_data = con_data.drop(['tagline', 'homepage'], axis=1)
```

```
In [9]: # Dropped the rows with missing values since they are few
con_data = con_data.dropna(axis = 0, how ='any')
```

```
In [14]: con_data.columns
```

```
Out[14]: Index(['title', 'genres', 'id', 'keywords', 'original_language', 'overview',
               'production_companies', 'production_countries', 'release_date',
               'runtime', 'spoken_languages', 'status', 'budget', 'revenue'],
              dtype='object')
```

```
In [15]: # Calculate Correlation
con_data.describe().T
```

```
Out[15]:
```

	count	mean	std	min	25%	50%	75%	max
id	4799.0	5.689992e+04	8.823650e+04	5.0	9012.5	14623.0	58461.5	4.470270e+05
runtime	4799.0	1.069031e+02	2.256131e+01	0.0	94.0	103.0	118.0	3.380000e+02
budget	4799.0	2.906593e+07	4.073251e+07	0.0	800000.0	15000000.0	40000000.0	3.800000e+08
revenue	4799.0	8.232920e+07	1.629076e+08	0.0	0.0	19184015.0	92956519.0	2.787965e+09

```
In [16]: con_data['release_date']
```

```
Out[16]: 0      12/10/2009
          1      5/19/2007
          2     10/26/2015
          3      7/16/2012
          4      3/7/2012
          5      5/1/2007
          6     11/24/2010
          7      4/22/2015
          8      7/7/2009
          9      3/23/2016
         10      6/28/2006
         11     10/30/2008
         12      6/20/2006
         13      7/3/2013
         14      6/12/2013
         15      5/15/2008
         16      4/25/2012
         17      5/14/2011
         18      5/23/2012
         19     12/10/2014
         20      6/27/2012
         21      5/12/2010
         22     12/11/2013
         23      12/4/2007
         24     12/14/2005
         25     11/18/1997
         26      4/27/2016
         27      4/11/2012
         28      6/9/2015
         29     10/25/2012
          ...
        4773     9/13/1994
        4774      1/1/1971
        4775     9/20/2002
        4776     1/19/1997
        4777     1/15/2002
        4778     8/14/2009
        4779     7/27/1990
        4780     10/2/2015
        4781     2/14/2013
        4782      1/1/2003
        4783     1/16/2015
        4784     1/17/2005
        4785      9/5/2014
        4786     3/14/2009
        4787     10/26/2011
        4788     3/12/1972
        4789      9/1/2004
        4790      9/8/2000
        4791      1/1/2007
        4792     11/6/1997
        4793     4/11/2004
        4794     1/20/2012
        4795      9/9/1995
        4796     10/8/2004
        4797     3/12/2005
        4798      9/4/1992
        4799     12/26/2011
        4800     10/13/2013
        4801      5/3/2012
        4802      8/5/2005
Name: release_date, Length: 4799, dtype: object
```



```
In [10]: # correcting the dates columns in a standard format
con_data[['release_month', 'release_day', 'release_year']] = con_data['release_date'].str.split('/', expand=True).replace(np.nan, -1).astype(int)

#getting the month year and day using the string split function and the / as a delimiter; eg: 5/25/2015 -> month 5/ day 25 / year 2015
con_data.loc[(con_data['release_year'] <= 19) & (con_data['release_year'] < 100),
"release_year"] += 2000

## some rows have 4 digits for the year instead of 2, so the release year < 100 and > 100 is checking that
con_data.loc[(con_data['release_year'] > 19) & (con_data['release_year'] < 100),
"release_year"] += 1900

releaseDate = pd.to_datetime(con_data['release_date'])

con_data.head()
```

Out [10]:

	title	genres	id	keywords	original_language	overview	production_companies	production_
0	Avatar	[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 80, 'name': 'Science Fiction'}, {'id': 27, 'name': 'Thriller'}]	19995	[{'id': 1463, 'name': 'culture clash'}, {'id': 1464, 'name': 'marine'}]	en	In the 22nd century, a paraplegic Marine is dispatched to the moon Pandora on a unique mission, but becomes torn between following orders and protecting those who have become his family.	[{'name': 'Ingenious Film Partners', 'id': 289}, {'name': 'Lightstorm Entertainment', 'id': 1}],	[{'iso_3166_1': 'US', 'name': 'United States'}]
1	Pirates of the Caribbean: At World's End	[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 27, 'name': 'Thriller'}]	285	[{'id': 270, 'name': 'ocean'}, {'id': 726, 'name': 'pirate'}, {'id': 80, 'name': 'science fiction'}, {'id': 1464, 'name': 'marine'}]	en	Captain Barbossa, long believed to be dead, has returned to haunt Will Turner and the rest of the crew of the Flying Dutchman.	[{'name': 'Walt Disney Pictures', 'id': 2}, {'name': 'Piracy Pictures', 'id': 1}],	[{'iso_3166_1': 'US', 'name': 'United States'}]
2	Spectre	[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 80, 'name': 'Science Fiction'}, {'id': 27, 'name': 'Thriller'}]	206647	[{'id': 470, 'name': 'spy'}, {'id': 818, 'name': 'mystery'}, {'id': 1464, 'name': 'marine'}]	en	A cryptic message from Bond's past sends him on a new mission.	[{'name': 'Columbia Pictures', 'id': 5}, {'name': 'United Artists', 'id': 1}],	[{'iso_3166_1': 'US', 'name': 'United States'}]
3	The Dark Knight Rises	[{'id': 28, 'name': 'Action'}, {'id': 80, 'name': 'Science Fiction'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 27, 'name': 'Thriller'}]	49026	[{'id': 849, 'name': 'dc comics'}, {'id': 853, 'name': 'superhero'}, {'id': 1464, 'name': 'marine'}]	en	Following the death of District Attorney Harvey Dent, Batman deduces that the only person who could have framed him for Dent's death is the mysterious figure known as the Joker.	[{'name': 'Legendary Pictures', 'id': 923}, {'name': 'Warner Bros. Entertainment', 'id': 1}],	[{'iso_3166_1': 'US', 'name': 'United States'}]
4	John Carter	[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}, {'id': 80, 'name': 'Science Fiction'}, {'id': 27, 'name': 'Thriller'}]	49529	[{'id': 818, 'name': 'based on novel'}, {'id': 1464, 'name': 'marine'}]	en	John Carter is a war-weary, former military captain, plagued by traumatic experiences from a long and troubled past in the American Civil War.	[{'name': 'Walt Disney Pictures', 'id': 2}],	[{'iso_3166_1': 'US', 'name': 'United States'}]

```
In [11]: data = mvrevenue
data[['release_month', 'release_day', 'release_year']] = data['release_date'].str.split(
    '/', expand=True).replace(np.nan, -1).astype(int) #getting the month year and day u
sing the string split function and the / as a delimiter; eg: 5/25/2015 -> month 5/
day 25 / year 2015
data.loc[ (data['release_year'] <= 19) & (data['release_year'] < 100), "release_yea
r"] += 2000 ## some rows have 4 digits for the year instead of 2, so the release ye
ar < 100 and > 100 is checking that
data.loc[ (data['release_year'] > 19) & (data['release_year'] < 100), "release_yea
r"] += 1900

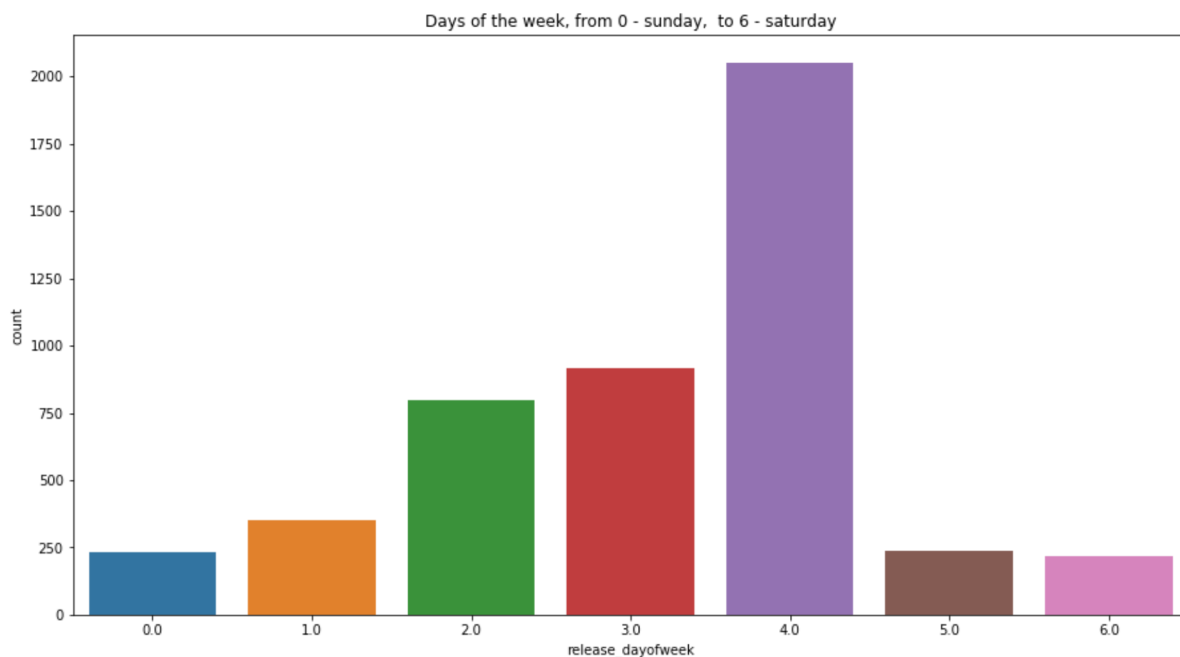
releaseDate = pd.to_datetime(data['release_date']) #using the pandas to_datetime fu
nction to format the data, get a Series, and store it in a variable that is gonna
be used later to get the day of week and quarter
data['release_dayofweek'] = releaseDate.dt.dayofweek
data['release_quarter'] = releaseDate.dt.quarter
```

```
In [12]: plt.figure(figsize=(15, 8))
sns.countplot(data['release_dayofweek'])
plt.title('Days of the week, from 0 - sunday, to 6 - saturday')
```

Out[12]: <Figure size 1080x576 with 0 Axes>

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1f71f9baa90>

Out[12]: Text(0.5, 1.0, 'Days of the week, from 0 - sunday, to 6 - saturday')




```
In [13]: plt.figure(figsize=(15, 8))  
sns.countplot(data['release_month']); plt.title('Distribution of movies by month')
```

Out[13]: <Figure size 1080x576 with 0 Axes>

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1f72588d390>

Out[13]: Text(0.5, 1.0, 'Distribution of movies by month')




```
In [23]: plt.figure(figsize=(15, 8))
sns.countplot(data['release_day']); plt.title('Distribution of movies by days in a month')
```

Out[23]: <Figure size 1080x576 with 0 Axes>

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x23682dfedd8>

Out[23]: Text(0.5, 1.0, 'Distribution of movies by days in a month')



In []:

In []:

Data Description Report

Issues found:

- dates column
- text to dict column

Exploratory Data Analysis

This section handles the graphs and plots for data exploration.

checking for outliers & anomalies in runtime, budget and revenue

- Univariate visualization

Univariate analysis looks at one feature at a time. When we analyze a feature independently, we are usually mostly interested in the distribution of its values and ignore other features in the dataset.

Below, we will consider different statistical types of features and the corresponding tools for their individual visual analysis.

- Quantitative features

Quantitative features take on ordered numerical values. Those values can be discrete, like integers, or continuous, like real numbers, and usually express a count or a measurement.


- Frequency distributions and class distributions

Plotting distplot to check the distribution.

```
In [74]: plt.figure(figsize=(16,10))  
sns.distplot(con_data['budget'])
```

```
Out[74]: <Figure size 1152x720 with 0 Axes>
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x22ec016bb38>
```



Analysis

the above plots show the 'budget' contain outliers, with distribution as positively skewed and a thin kurtosis.

```
In [14]: # creating gen dataframe for genres to be compared against revenue, budget, runtime
         and status
gen = con_data.loc[con_data['genres'].str.len()==1][
    ['genres', 'revenue', 'budget', 'runtime', 'status']].reset_index(drop = True)
gen['genres'] = gen.genres.apply(lambda x :x[0]['name'])

#gen['genres'] = mvrevenue['genres'].str.extract('([A-Z]\w{0,})', expand=True )

gen.head()
```

Out [14]:

	genres	revenue	budget	runtime	status
0	Action	1506249360	190000000	137.0	Released
1	Science Fiction	543934787	178000000	144.0	Released
2	Fantasy	299370084	170000000	113.0	Released
3	Adventure	301000000	150000000	99.0	Released
4	Comedy	202026112	100000000	90.0	Released

```
In [73]: #del mv0

#mv0 = mvrevenue.copy()
#mv0 = mv0['genres'].str.extract('^[A-Za-z]+', expand=True )

#mv0 = mvrevenue['genres'].str.extract('([A-Z]\w{0,})', expand=True )

#mv0
```

```
In [15]: genres = gen.groupby(gen.genres).agg('mean')
genres
```

Out [15]:

	revenue	budget	runtime
genres			
Action	1.227577e+08	3.971429e+07	107.904762
Adventure	9.679045e+07	4.853846e+07	110.461538
Animation	1.231166e+08	4.950000e+07	93.000000
Comedy	5.707003e+07	2.031676e+07	96.184397
Crime	0.000000e+00	2.500000e+06	110.000000
Documentary	5.566236e+06	1.735243e+06	92.909091
Drama	3.436563e+07	1.298058e+07	114.425474
Family	0.000000e+00	0.000000e+00	74.000000
Fantasy	9.671589e+07	6.483333e+07	112.833333
History	2.652751e+07	1.800000e+07	120.000000
Horror	3.999399e+07	8.085078e+06	89.765625
Music	1.423834e+07	1.600000e+07	98.000000
Romance	6.508158e+07	8.340000e+06	96.000000
Science Fiction	7.420446e+07	3.829107e+07	110.142857
Thriller	2.070070e+07	1.486957e+07	107.086957
War	6.190232e+07	4.300000e+07	126.500000
Western	2.279019e+07	7.017988e+06	119.277778

```
In [20]: plt.figure(figsize=(15,10))
plt.subplot(2,2,3)
sns.barplot(genres['revenue'],genres.index)

plt.subplot(2,2,2)
sns.barplot(genres['runtime'],genres.index)

plt.subplot(2,2,1)
sns.barplot(genres['budget'],genres.index)
```

Out[20]: <Figure size 1080x720 with 0 Axes>

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14024ff60>



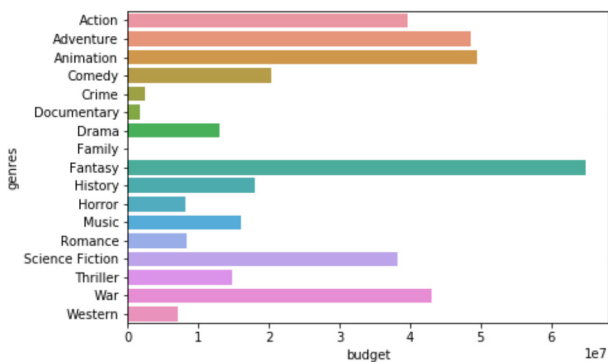
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14024ff60>

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c140279828>

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c140279828>

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14036e7f0>

Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14036e7f0>




```
Out[34]: original_language
af      1
ar      2
cn     12
cs      2
da      7
de     26
el      1
en    4503
es     32
fa      4
fr     70
he      3
hi     19
hu      1
id      2
is      1
it     13
ja     16
ko     11
ky      1
nb      1
nl      4
no      1
pl      1
ps      1
pt      9
ro      2
ru     11
sl      1
sv      5
ta      2
te      1
th      3
tr      1
vi      1
xx      1
zh     27
dtype: int64
```

DF backup before further modification

Out [16] :

	title	genres	id	keywords	original_language	overview	production_companies	production_countries
0	Avatar	[{'id': 28, 'name': 'Action'}, {'id': 12, 'nam...	19995	[{'id': 1463, 'name': 'culture clash'}, {'id':	en	In the 22nd century, a paraplegic Marine is di...	[{'name': 'Ingenious Film Partners', 'id': 289...}	[{'iso_3166_1': 'US', 'name': 'United States', 'o...

```
In [17]: element = con_data2.iloc[0]['genres']  
         element
```

```
Out[17]: [{'id': 28, 'name': 'Action'},  
          {'id': 12, 'name': 'Adventure'},  
          {'id': 14, 'name': 'Fantasy'},  
          {'id': 878, 'name': 'Science Fiction'}]
```

Converting the dictionary columns to extract the values

```
In [18]: def parse_dict(raw_dict):  
         return [d['name'] for d in raw_dict ]
```

```
In [40]: con_data2
```

	title	genres	id	keywords	original_language	overview	production_compa
0	Avatar	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	19995	[[{'id': 1463, 'name': 'culture clash'}, {'id': 12, 'name': 'culture clash'}], {'id': 12, 'name': 'culture clash'}	en	In the 22nd century, a paraplegic Marine is di...	[[{'name': 'Ingenious Partners', 'id': 2}], {'name': 'Ingenious Partners', 'id': 2}]
1	Pirates of the Caribbean: At World's End	[[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}], {'id': 14, 'name': 'Fantasy'}	285	[[{'id': 270, 'name': 'ocean'}, {'id': 726, 'name': 'pirates'}], {'id': 726, 'name': 'pirates'}	en	Captain Barbossa, long believed to be dead, ha...	[[{'name': 'Walt Disney Pictures', 'id': 2}], {'name': 'Walt Disney Pictures', 'id': 2}]
2	Spectre	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	206647	[[{'id': 470, 'name': 'spy'}, {'id': 818, 'name': 'mystery'}], {'id': 818, 'name': 'mystery'}	en	A cryptic message from Bond's past sends him o...	[[{'name': 'Columbia Pictures', 'id': 5}], {'name': 'Columbia Pictures', 'id': 5}]
3	The Dark Knight Rises	[[{'id': 28, 'name': 'Action'}, {'id': 80, 'name': 'Fantasy'}], {'id': 80, 'name': 'Fantasy'}	49026	[[{'id': 849, 'name': 'dc comics'}, {'id': 853, 'name': 'superhero'}], {'id': 853, 'name': 'superhero'}	en	Following the death of District Attorney Harvey Dent, Batman seeks out a new ally to bring justice to Gotham City...	[[{'name': 'Legendary Pictures', 'id': 923}], {'name': 'Legendary Pictures', 'id': 923}]
4	John Carter	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	49529	[[{'id': 818, 'name': 'based on novel'}, {'id': 12, 'name': 'fantasy'}], {'id': 12, 'name': 'fantasy'}	en	John Carter is a war-weary, former military captain who mysteriously finds himself transported to a strange, war-torn planet...	[[{'name': 'Walt Disney Pictures', 'id': 2}], {'name': 'Walt Disney Pictures', 'id': 2}]
5	Spider-Man 3	[[{'id': 14, 'name': 'Fantasy'}, {'id': 28, 'name': 'Action'}], {'id': 28, 'name': 'Action'}	559	[[{'id': 851, 'name': 'dual identity'}, {'id': 12, 'name': 'fantasy'}], {'id': 12, 'name': 'fantasy'}	en	The seemingly invincible Spider-Man goes up against his greatest foe yet...	[[{'name': 'Columbia Pictures', 'id': 5}], {'name': 'Columbia Pictures', 'id': 5}]
6	Tangled	[[{'id': 16, 'name': 'Animation'}, {'id': 10751, 'name': 'Fantasy'}], {'id': 10751, 'name': 'Fantasy'}	38757	[[{'id': 1562, 'name': 'hostage'}, {'id': 2343, 'name': 'romance'}], {'id': 2343, 'name': 'romance'}	en	When the kingdom's most wanted-and most charming-Prince Eric is kidnapped by a vengeful sorceress...	[[{'name': 'Walt Disney Pictures', 'id': 2}], {'name': 'Walt Disney Pictures', 'id': 2}]
7	Avengers: Age of Ultron	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	99861	[[{'id': 8828, 'name': 'marvel comic'}, {'id': 12, 'name': 'fantasy'}], {'id': 12, 'name': 'fantasy'}	en	When Tony Stark tries to jumpstart a dormant prototype, he unleashes a powerful, intelligent, and unstoppable force...	[[{'name': 'Marvel Studios', 'id': 4}], {'name': 'Marvel Studios', 'id': 4}]
8	Harry Potter and the Half-Blood Prince	[[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}], {'id': 14, 'name': 'Fantasy'}	767	[[{'id': 616, 'name': 'witch'}, {'id': 2343, 'name': 'magic'}], {'id': 2343, 'name': 'magic'}	en	As Harry begins his sixth year at Hogwarts, he discovers an unbreakable bond between himself and a powerful, dark wizard...	[[{'name': 'Warner Bros. Pictures', 'id': 6194}], {'name': 'Warner Bros. Pictures', 'id': 6194}]
9	Batman v Superman: Dawn of Justice	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	209112	[[{'id': 849, 'name': 'dc comics'}, {'id': 7002, 'name': 'superhero'}], {'id': 7002, 'name': 'superhero'}	en	Fearing the actions of a god-like Super Hero have provoked a head-on clash with the forces of government...	[[{'name': 'DC Comics', 'id': 429}], {'name': 'DC Comics', 'id': 429}]
10	Superman Returns	[[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}], {'id': 14, 'name': 'Fantasy'}	1452	[[{'id': 83, 'name': 'saving the world'}, {'id': 12, 'name': 'fantasy'}], {'id': 12, 'name': 'fantasy'}	en	Superman returns to discover his 5-year absence has left the world a different place...	[[{'name': 'DC Comics', 'id': 429}], {'name': 'DC Comics', 'id': 429}]
11	Quantum of Solace	[[{'id': 12, 'name': 'Adventure'}, {'id': 28, 'name': 'Action'}], {'id': 28, 'name': 'Action'}	10764	[[{'id': 627, 'name': 'killing'}, {'id': 1568, 'name': 'romance'}], {'id': 1568, 'name': 'romance'}	en	Quantum of Solace continues the adventures of James Bond as he uncovers a conspiracy that leads to a new and more powerful enemy...	[[{'name': 'Sony Pictures Productions', 'id': 75}], {'name': 'Sony Pictures Productions', 'id': 75}]
12	Pirates of the Caribbean: Dead Man's Chest	[[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}], {'id': 14, 'name': 'Fantasy'}	58	[[{'id': 616, 'name': 'witch'}, {'id': 663, 'name': 'pirates'}], {'id': 663, 'name': 'pirates'}	en	Captain Jack Sparrow works his way out of a bind when he uncovers an ancient treasure of the legendary Dutch Golden Age captain...	[[{'name': 'Walt Disney Pictures', 'id': 2}], {'name': 'Walt Disney Pictures', 'id': 2}]
13	The Lone Ranger	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	57201	[[{'id': 1556, 'name': 'texas'}, {'id': 2673, 'name': 'western'}], {'id': 2673, 'name': 'western'}	en	The Texas Rangers chase down a gang of outlaws in a high-stakes battle for justice...	[[{'name': 'Walt Disney Pictures', 'id': 2}], {'name': 'Walt Disney Pictures', 'id': 2}]
14	Man of Steel	[[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Fantasy'}], {'id': 12, 'name': 'Fantasy'}	49521	[[{'id': 83, 'name': 'saving the world'}, {'id': 12, 'name': 'fantasy'}], {'id': 12, 'name': 'fantasy'}	en	A young boy learns that he has extraordinary powers as he discovers his Kryptonian heritage...	[[{'name': 'Legendary Pictures', 'id': 923}], {'name': 'Legendary Pictures', 'id': 923}]

```
In [19]: # One-Hot-Encoding for all nominal data

df = pd.get_dummies(gen)

df.head()
```

Out [19]:


	revenue	budget	runtime	genres_Action	genres_Adventure	genres_Animation	genres_Comedy	gen
0	1506249360	190000000	137.0	1	0	0	0	
1	543934787	178000000	144.0	0	0	0	0	
2	299370084	170000000	113.0	0	0	0	0	
3	301000000	150000000	99.0	0	1	0	0	
4	202026112	100000000	90.0	0	0	0	1	

5 rows × 23 columns

```
In [42]: sns.set(rc={'figure.figsize': (11.7, 8.27)})

sns.distplot(df['revenue'])
```


Out [42]: <matplotlib.axes._subplots.AxesSubplot at 0x1a55a0dd2e8>



```
In [43]: sns.set(rc={'figure.figsize': (11.7, 8.27)})
```

```
sns.distplot(df['runtime'])
```


```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x1a55a1f3278>
```



```
In [44]: sns.set(rc={'figure.figsize': (11.7, 8.27)})
```

```
sns.distplot(df['budget'])
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x1a55a2c5518>
```



```
In [20]: # Skewness value  
con_data.skew()
```

```
Out[20]: id                2.071986  
runtime            0.739876  
budget             2.436115  
revenue            4.443129  
release_month      -0.153424  
release_day         0.022664  
release_year       -2.170769  
dtype: float64
```

Modeling

```
In [21]: pd.set_option('display.max_columns', None)  
pd.set_option('display.expand_frame_repr', False)  
pd.set_option('max_colwidth', -1)
```

```
In [22]: df.head()
```

```
Out [22]:
```

	revenue	budget	runtime	genres_Action	genres_Adventure	genres_Animation	genres_Comedy	gen
0	1506249360	1900000000	137.0	1	0	0	0	
1	543934787	1780000000	144.0	0	0	0	0	
2	299370084	1700000000	113.0	0	0	0	0	
3	301000000	1500000000	99.0	0	1	0	0	
4	202026112	1000000000	90.0	0	0	0	1	

```
In [27]: df.columns
```

```
Out [27]: Index(['revenue', 'budget', 'runtime', 'genres_Action', 'genres_Adventure',
                'genres_Animation', 'genres_Comedy', 'genres_Crime',
                'genres_Documentary', 'genres_Drama', 'genres_Family', 'genres_Fantasy',
                'genres_History', 'genres_Horror', 'genres_Music', 'genres_Romance',
                'genres_Science Fiction', 'genres_Thriller', 'genres_War',
                'genres_Western', 'status_Post Production', 'status_Released',
                'status_Rumored'],
                dtype='object')
```

Split the data into training set and testing set using train_test_split

using scikit learn split the data-set

```
In [23]: df = df[['budget', 'runtime', 'genres_Action', 'genres_Adventure',
                'genres_Animation', 'genres_Comedy', 'genres_Crime',
                'genres_Documentary', 'genres_Drama', 'genres_Family', 'genres_Fantasy',
                'genres_Horror', 'genres_Science Fiction', 'genres_Thriller',
                'genres_Western', 'status_Released', 'revenue']]
```

```
In [24]: # Split-out validation dataset

X = df.drop('revenue', axis = 1)
Y = df.revenue
validation_size = 0.20
seed = 7
X_train, X_validation, Y_train, Y_validation = train_test_split(X, Y, test_size=validation_size, random_state=seed)
```



```
In [30]: print(X_train)
         print(X_validation)
         print(Y_train)
         print(Y_validation)
```

	budget	runtime	genres_Action	genres_Adventure	genres_Animation	genres_Comedy	genres_Crime	genres_Documentary	genres_Drama	genres_Family	genres_Fantasy	genres_Horror	genres_Science Fiction	genres_Thriller	genres_Western	status_Released
264	35000000	114.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
22	79000000	91.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
815	400000	95.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
888	27000	92.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
835	0	95.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
232	27000000	113.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
880	50000	111.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
270	20000000	127.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
669	35000000	94.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
316	18000000	111.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
465	11000000	98.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
347	26000000	125.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
142	35000000	91.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
50	58000000	103.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
709	1500000	111.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
885	0	89.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
445	6000000	90.0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
558	5000000	104.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
750	0	90.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
81	0	0.0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0

Running the Linear Regression Model

```
In [26]: from sklearn.linear_model import LinearRegression
reg = LinearRegression()
def rmsle(y, y0): return np.sqrt(np.mean(np.square(np.log1p(y) - np.log1p(y0))))
model = reg.fit(X, Y)
y_pred = reg.predict(X)
rmsle = rmsle(y_pred, Y)
print("The linear model has intercept : {}, and coefficients : {}, and the rmsle is
{} ".format(model.intercept_, model.coef_, rmsle) )
```

The linear model has intercept : -14644620.058988929, and coefficients : [2.118
93237e+00 2.61006990e+05 2.78379824e+07 -1.74944018e+07
1.13518567e+07 6.30172914e+06 -1.66119971e+07 -5.00613574e+06
-5.61654625e+06 -1.91841449e+06 -5.27157457e+07 1.68289014e+07
-1.82836835e+07 -2.13612411e+07 -5.81668017e+06 -2.75148271e+06], and the rmsle
is 9.36410874773492

```
In [27]: # Build and fit linear regression model
reg_lm = LinearRegression(normalize=True)
reg_lm.fit(X_train, Y_train)

# Calculate the predictions on the test set
pred = reg_lm.predict(X_validation)
```

Out[27]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=True)

In [28]: pred

```
Out[28]: array([-2.19615413e+06,  1.17301722e+08,  1.60152199e+07,  5.17659998e+06,
  4.75559258e+07, -4.62763531e+06,  1.10256610e+07,  4.09233613e+07,
  3.00847413e+07,  3.63919353e+07,  1.65019210e+06,  3.50653250e+07,
  5.67876503e+07,  1.24129311e+07,  4.40054836e+05, -1.38925141e+07,
 -5.13377653e+05,  1.33745176e+07,  3.95415235e+07,  2.77168851e+07,
  2.60362459e+06,  1.07141608e+08,  2.82808808e+07, -2.31106618e+07,
  6.85688571e+07,  2.21896516e+07,  4.64732552e+07,  1.04095993e+08,
  1.23267064e+08,  4.46069634e+07,  1.13076000e+07,  7.81273325e+07,
 -2.72981013e+05,  4.48208905e+05,  5.56320122e+07,  7.11336784e+07,
  7.02052020e+07,  7.59018404e+07, -1.23456757e+06,  1.38553109e+07,
  2.04842147e+07,  2.77263359e+06,  5.18494580e+06,  3.65717899e+07,
  3.91229167e+07,  1.07701943e+08,  6.00144998e+07, -1.64157930e+07,
  5.01360861e+07,  2.07852125e+08,  5.01574505e+06,  4.40054836e+05,
  1.89058874e+06,  8.07910864e+07,  1.70889233e+08,  7.10370263e+07,
  1.30520506e+08,  6.77681255e+07,  1.43361042e+07,  2.45709576e+07,
  9.33473049e+06,  8.92271203e+07,  2.53183898e+07,  6.69036746e+06,
 -1.54144356e+07,  1.01588371e+08,  2.08041531e+07,  8.62169465e+06,
 -1.47496421e+06,  6.05146839e+07,  1.34510597e+08,  1.01131692e+08,
  5.02571702e+07,  5.76525847e+07,  3.60548865e+07,  5.01574505e+06,
  1.03134406e+08,  9.15359678e+07,  2.84402123e+06,  3.46304071e+07,
  1.72208639e+07,  1.42809185e+07,  7.59481855e+07,  1.00490043e+08,
  1.20982411e+08,  1.26533277e+07,  2.52523766e+07,  1.02131753e+07,
  1.48168975e+07, -3.45417082e+06,  1.84833891e+07,  1.52334880e+07,
  2.60714367e+07,  1.48168975e+07,  3.48529033e+07,  2.53782911e+07,
  1.36149143e+07, -2.81135082e+05,  5.48889091e+07,  5.44697999e+07,
  5.97733161e+06,  1.50572941e+07,  5.89756519e+07,  2.53782635e+07,
  1.16917412e+07,  5.53148109e+07,  1.57784840e+07,  1.25493541e+07,
  5.73693497e+06,  1.36149143e+07,  4.52315011e+07,  6.36014497e+07,
  3.80560779e+06, -2.21765900e+07,  2.30716964e+07,  1.55551978e+07,
 -9.94170932e+05,  1.76184547e+07,  8.02136409e+07,  2.37138202e+06,
 -3.90644539e+06,  6.16200150e+07,  1.12167277e+08,  3.80560779e+06,
  9.15762410e+07,  1.16591552e+07,  2.37892256e+07,  5.82366044e+07,
  3.98216910e+07,  4.09843789e+06,  2.12886071e+07,  2.37138202e+06,
 -1.23456757e+06,  1.23762017e+05,  2.41697060e+07,  7.86275166e+07,
  6.15779637e+07,  9.64828422e+07,  6.10815932e+07,  5.32856743e+07,
 -1.82196536e+07,  1.22208826e+08, -5.13377653e+05,  1.77486129e+07,
  2.89164161e+07,  2.95795162e+07,  3.81376186e+06,  4.60078009e+07,
  8.09402088e+07,  1.31700556e+07,  1.18281642e+07,  5.73693497e+06,
  6.43610302e+07,  1.68358040e+06,  1.28937243e+07,  1.88243467e+06,
  2.37755312e+07,  2.82808808e+07,  7.52218016e+07,  3.80747118e+07,
  5.73693497e+06,  1.34313989e+07,  2.19318808e+08,  7.00754397e+07,
  2.16861979e+07, -4.83508493e+06, -1.47496421e+06,  1.74178289e+07,
  3.24310793e+07,  1.32664685e+08, -1.58156872e+07,  5.80140654e+07,
  3.65356797e+07,  4.05415850e+06,  5.12982984e+07,  9.63617581e+07,
  1.32106630e+08,  1.83069438e+07,  1.88645418e+07,  4.04624155e+07])
```

```
In [29]: print('*****Y Shape*****')
Y_validation.shape

print('*****Pred Shape*****')
pred.shape
```

*****Y Shape*****

Out[29]: (180,)

*****Pred Shape*****

Out[29]: (180,)

Assess Model

```
In [30]: # Evaluate the performance using the RMSE
from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(Y_validation, pred))
print('RMSE: {:.3f}'.format(rmse))
```


RMSE: 64925743.344

```
In [31]: # r2_score is between 0 & 1. 1 is best fit & 0 is worst fit
from sklearn.metrics import r2_score
print("r2_score", r2_score(Y_validation, pred))
```

r2_score 0.19900571473920703

```
In [32]: # visualize the relationship between the features and the response using scatterplots
sns.pairplot(mvrevenue, x_vars=['budget', 'release_year', 'runtime'], y_vars='revenue', size=7, aspect=0.7)
sns.pairplot(df, x_vars=[ 'budget', 'runtime'], y_vars='revenue', size=7, aspect=0.7)
```

Out [32]: <seaborn.axisgrid.PairGrid at 0x1f727d5c5c0>



```
In [34]: # create X and y
feature_cols = ['budget']
X = df[feature_cols]
y = df.revenue

# instantiate and fit
lm2 = LinearRegression()
lm2.fit(X, y)

# print the coefficients
print(lm2.intercept_)
print(lm2.coef_)
```

```
Out[34]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

8544030.1107831
[2.17718383]
```

Linear Regression Explain

$y = \beta_1 \cdot x_1$, here β_1 is the coefficient, x_1 is the budget and y is the revenue.

Customer wanted to predict the budget allocation, which can be calculated as follows

Interpreting Model Coefficients

Interpreting the budget coefficient (β_1)


So, the budget unit can be calculated as $x_1 = y/\beta_1$,
(8544030.1107831/ 2.17718383) = 3,924.68075

```
In [ ]:
```

Plotting the Least Squares Line

```
In [35]: sns.pairplot(mvrevenue, x_vars=['budget', 'runtime'], y_vars='revenue', size=7, aspect=0.7, kind='reg')
```

```
Out[35]: <seaborn.axisgrid.PairGrid at 0x1f728606748>
```



```
In [36]: # create X and y
feature_cols = ['budget', 'runtime']
X = df[feature_cols]
y = df.revenue

# # instantiate and fit
lm2 = LinearRegression()
lm2.fit(X, y)

# # print the coefficients
print(lm2.intercept_)
print(lm2.coef_)
```

```
Out[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
-5309987.732538827
[2.14442863e+00 1.37603060e+05]
```

```
In [37]: # # instantiate and fit
lm2 = LinearRegression()
lm2.fit(X, y)
```

```
Out[37]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

Predict with Linear Regression

```
In [38]: X_train, X_validation, Y_train, Y_validation = train_test_split(X, Y, test_size=val
         idation_size, random_state=seed)
```

```
In [39]: # Remove table meta data, column names to use values for prediction.

# train_x = df_train_x.values
# train_y = df_train_y.values
# test_x = df_test_x.values
# X_train = X_train.values
# Y_train = Y_train.values
# Calculate the coefficients of the linear regression
reg = LinearRegression().fit(X_train, Y_train)

# Using linear regression model on the prepared test data
Y_validation = reg.predict(X_validation)

# Accuracy
print('Accuracy Linear Regression:', reg.score(X_validation, Y_validation))
```

Accuracy Linear Regression: 1.0

Predicting with XGBOOST

```
In [40]: import xgboost as xgb

xgb_model = xgb.XGBRegressor(objective="reg:squarederror", random_state=7)

xgb_model.fit(X_train, Y_train)

Y_validation = xgb_model.predict(X_validation)

print('Accuracy XGB:', xgb_model.score(X_train, Y_train))
```

```
Out[40]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, gamma=0,
                    importance_type='gain', learning_rate=0.1, max_delta_step=0,
                    max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                    n_jobs=1, nthread=None, objective='reg:squarederror',
                    random_state=7, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                    seed=None, silent=None, subsample=1, verbosity=1)
```

Accuracy XGB: 0.7066468166703885

Trying with other algorithms

```
In [41]: # Test options and evaluation metric
num_folds = 10
num_instances = len(X_train)
seed = 7

# Listing the possible scoring matrix

r2_score = 'r2'
## metrics.r2_score

# Initiating the score matrix
scoring = r2_score
```



```
In [44]: random_seed = 12
outcome = []
model_names = []
models = []
models.append(('LR', LinearRegression()))
models.append(('LASSO', Lasso()))
models.append(('EN', ElasticNet()))
models.append(('KNN', KNeighborsRegressor()))
models.append(('CART', DecisionTreeRegressor()))
models.append(('SVR', SVR()))
models.append(('XGB', xgb.XGBRegressor(objective="reg:squarederror")))
```

```
In [45]: # evaluate each model in turn
results = []
names = []
for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state=random_seed)
    cv_results = model_selection.cross_val_score(model, X_train, Y_train, cv=kfold,
scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

```
LR: 0.347419 (0.122490)
LASSO: 0.347419 (0.122490)
EN: 0.347420 (0.122491)
KNN: 0.263239 (0.159443)
CART: -0.094588 (0.291035)
SVR: -0.222887 (0.094258)
XGB: 0.334331 (0.146317)
```

```
In [46]: # Validating LR Score on Testing Set
LR_model = LinearRegression()
LR_model.fit(X_train, Y_train)
Y_validation = LR_model.predict(X_validation)
print('Accuracy LR:', LR_model.score(X_train, Y_train))

# Validating LASSO Score on Testing Set
LASSO_model = Lasso()
LASSO_model.fit(X_train, Y_train)
Y_validation = LR_model.predict(X_validation)
print('Accuracy LASSO:', LASSO_model.score(X_train, Y_train))

# Validating EN Score on Testing Set
EN_model = ElasticNet()
EN_model.fit(X_train, Y_train)
Y_validation = EN_model.predict(X_validation)
print('Accuracy EN:', EN_model.score(X_train, Y_train))

# Validating KNN Score on Testing Set
KNN_model = KNeighborsRegressor()
KNN_model.fit(X_train, Y_train)
Y_validation = KNN_model.predict(X_validation)
print('Accuracy KNN:', KNN_model.score(X_train, Y_train))

# Validating CART Score on Testing Set
CART_model = DecisionTreeRegressor()
CART_model.fit(X_train, Y_train)
Y_validation = CART_model.predict(X_validation)
print('Accuracy CART:', CART_model.score(X_train, Y_train))

# Validating SVR Score on Testing Set
SVR_model = SVR()
SVR_model.fit(X_train, Y_train)
Y_validation = SVR_model.predict(X_validation)
print('Accuracy SVR:', SVR_model.score(X_train, Y_train))

# Validating XGB Score on Testing Set
xgb_model = xgb.XGBRegressor(objective="reg:squarederror")
xgb_model.fit(X_train, Y_train)
Y_validation = xgb_model.predict(X_validation)
print('Accuracy XGB:', xgb_model.score(X_train, Y_train))
```

```
Out[46]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Accuracy LR: 0.3829906441575285

Out[46]: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)

Accuracy LASSO: 0.3829906441575285

Out[46]: ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,
                    max_iter=1000, normalize=False, positive=False, precompute=False,
                    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

Accuracy EN: 0.38299064380781156

Out[46]: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                             weights='uniform')

Accuracy KNN: 0.5296165934823194

Out[46]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                               max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=1,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               presort=False, random_state=None, splitter='best')

Accuracy CART: 0.9807880333053189

Out[46]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
             gamma='auto_deprecated', kernel='rbf', max_iter=-1, shrinking=True,
             tol=0.001, verbose=False)

Accuracy SVR: -0.1582129454658463

Out[46]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0,
                      max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                      n_jobs=1, nthread=None, objective='reg:squarederror',
                      random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                      seed=None, silent=None, subsample=1, verbosity=1)


Accuracy XGB: 0.7066468166703885
```

```
In [47]: # Compare Algorithms
fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
plt.rcParams['figure.figsize'] = 14,10
```

```
Out[47]: Text(0.5, 0.98, 'Algorithm Comparison')
```

```
Out[47]: {'whiskers': [<matplotlib.lines.Line2D at 0x1f729cf62e8>,  
  <matplotlib.lines.Line2D at 0x1f729cf6630>,  
  <matplotlib.lines.Line2D at 0x1f729d01a20>,  
  <matplotlib.lines.Line2D at 0x1f729d01d68>,  
  <matplotlib.lines.Line2D at 0x1f729d0bef0>,  
  <matplotlib.lines.Line2D at 0x1f729d194e0>,  
  <matplotlib.lines.Line2D at 0x1f729d238d0>,  
  <matplotlib.lines.Line2D at 0x1f729d23c18>,  
  <matplotlib.lines.Line2D at 0x1f729d2fda0>,  
  <matplotlib.lines.Line2D at 0x1f729d3b390>,  
  <matplotlib.lines.Line2D at 0x1f729d45780>,  
  <matplotlib.lines.Line2D at 0x1f729d45ac8>,  
  <matplotlib.lines.Line2D at 0x1f729d52eb8>,  
  <matplotlib.lines.Line2D at 0x1f729d52f98>],  
  'caps': [<matplotlib.lines.Line2D at 0x1f729cf6978>,  
  <matplotlib.lines.Line2D at 0x1f729cf6cc0>,  
  <matplotlib.lines.Line2D at 0x1f729d01e48>,  
  <matplotlib.lines.Line2D at 0x1f729d0b438>,  
  <matplotlib.lines.Line2D at 0x1f729d19828>,  
  <matplotlib.lines.Line2D at 0x1f729d19b70>,  
  <matplotlib.lines.Line2D at 0x1f729d23f60>,  
  <matplotlib.lines.Line2D at 0x1f729d23ef0>,  
  <matplotlib.lines.Line2D at 0x1f729d3b6d8>,  
  <matplotlib.lines.Line2D at 0x1f729d3ba20>,  
  <matplotlib.lines.Line2D at 0x1f729d45e10>,  
  <matplotlib.lines.Line2D at 0x1f729d45ef0>,  
  <matplotlib.lines.Line2D at 0x1f729d5d588>,  
  <matplotlib.lines.Line2D at 0x1f729d5d8d0>],  
  'boxes': [<matplotlib.lines.Line2D at 0x1f729ca9438>,  
  <matplotlib.lines.Line2D at 0x1f729d016d8>,  
  <matplotlib.lines.Line2D at 0x1f729d0be10>,  
  <matplotlib.lines.Line2D at 0x1f729d23588>,  
  <matplotlib.lines.Line2D at 0x1f729d2fcc0>,  
  <matplotlib.lines.Line2D at 0x1f729d45438>,  
  <matplotlib.lines.Line2D at 0x1f729d52b70>],  
  'medians': [<matplotlib.lines.Line2D at 0x1f729cf6da0>,  
  <matplotlib.lines.Line2D at 0x1f729d0b780>,  
  <matplotlib.lines.Line2D at 0x1f729d19eb8>,  
  <matplotlib.lines.Line2D at 0x1f729d2f630>,  
  <matplotlib.lines.Line2D at 0x1f729d3bd68>,  
  <matplotlib.lines.Line2D at 0x1f729d524e0>,  
  <matplotlib.lines.Line2D at 0x1f729d5dc18>],  
  'fliers': [<matplotlib.lines.Line2D at 0x1f729d01390>,  
  <matplotlib.lines.Line2D at 0x1f729d0bac8>,  
  <matplotlib.lines.Line2D at 0x1f729d19f98>,  
  <matplotlib.lines.Line2D at 0x1f729d2f978>,  
  <matplotlib.lines.Line2D at 0x1f729d3be48>,  
  <matplotlib.lines.Line2D at 0x1f729d52828>,  
  <matplotlib.lines.Line2D at 0x1f729d5df60>],  
  'means': []}
```

```
Out[47]: [Text(0, 0, 'LR'),  
  Text(0, 0, 'LASSO'),  
  Text(0, 0, 'EN'),  
  Text(0, 0, 'KNN'),  
  Text(0, 0, 'CART'),  
  Text(0, 0, 'SVR'),  
  Text(0, 0, 'XGB')]
```



In []:

In []:

In []:

Review Process Review of Process

In []:

Determine Next Steps

In []:

List of Possible Actions Decision

In []:

Review Recommendations to Organization

In []:

In []: